

Accelerationism comes from the Marxist idea that communism comes after capitalism, and therefore, to get to communism, we have to accelerate the logics of capitalism. Rather than resisting corporate exploitation with robust unions, maybe we should implicate ourselves in corporate takeover of the world. Rather than feeling guilty about being a trashy consumer, maybe we should get out there and enjoy commodity fetishism. Accelerationism is a seductive idea. It says: it's okay to lay down the hermeneutics of suspicion, the nagging doubt that capitalism is out to get ya, to screw you over. It's all part of the plan. Just hang on tight, accelerate the process, and we'll be in communist utopia in no time. Never mind the question whether Marx was right; never mind the debate that well, Marx didn't *quite* mean that communism will just naturally happen to us inevitably; never minding those, nevertheless immediately one might worry: capitalism is racist. If we ought accelerate the logics of capitalism, doesn't this entail that we ought accelerate racism? But surely one would have to be out of one's mind to believe that.

There is a sizable contingent of people who believe precisely this, that we must accelerate racism; they call it "hyper-racism", and their method of praxis consists, mainly, of shitposting racist garbage on Twitter. And yes, they are out of their minds, they are batshit insane fascists. Anyone calling oneself an anti-racist should not be interacting with them. Anyway, we call them the Right-Accelerationists.

The comparatively sober antithesis to the Right-Accelerationists are, predictably, the Left-Accelerationists. The Left-Accelerationists are Marxists who think technology is kind of neat. They have a good knowledge of computer science, artificial intelligence, and the latest technological advancements, and believe technology is crucial for a proper accelerationism. But they don't often hang around the technical experts, and so when they claim some mathematical-ish statements have such and such philosophical implications, they are liable to go unchallenged. Emboldened, they try to philosophically tackle what seems like the holy grail of computer science, Artificial General Intelligence (AGI).

Intelligence and Spirit by Reza Negarestani is one such attempt. There are many things wrong with this book. Perhaps the book deserves a thorough excoriation, but that's not what I'm out to do here. So I will focus on one part which exemplifies the things that are wrong with the book. This is in chapter 7, "This I, or We or It, the Thing, Which Speaks (Language as Interaction as Computation)", the most important chapter of the book. Over a few pages, Mr. Negarestani attempts the argument that interaction is somehow more expressive than computation. Aside from the technical portion of the argument, which is incoherent, the motivation seems to be roughly that (1) interaction is analogous to the dialectic, and computation to deduction; (2) since we're all Hegelians here, we know the dialectic is superior to deduction; (3) so we know interaction is superior to computation. Rough as they are, I do not necessarily disagree with the second statement, and the third would logically follow from the first and second. It's just that the first statement fails to hold. Let me show exactly what is wrong with it.

Mr. Negarestani gets into his argument with a definition of the “classical Church-Turing computability thesis”:

According to the classical Church-Turing computability thesis, computation can be modelled on a mathematical function. [1]

Already we’re in trouble. This is *not* what the Church-Turing thesis says, “classical” or no. The thesis is controversial, and there are several variants, but the version pertaining to computability is usually put like this: “every effective computation can be carried out by a Turing machine”. I do not know how Mr. Negarestani equated “can be carried out by a Turing machine” with “can be modelled on a mathematical function”, but the seeds of his destruction are contained herein. A “mathematical function” is a far more ill-defined notion than a Turing machine. One way in which they differ is that a function is a mapping from input to output, and the output is assumed to exist. (Or not; the point is that it’s possible to quibble.) With a Turing machine, this assumption can no longer hold, because of the halting problem. While a mathematical function always outputs something, a Turing machine either outputs something, or runs forever. (It’s not possible to quibble.) This is a subtle point, and it is a point Mr. Negarestani is persistent on missing. In just another moment he says

[Classical] computation runs into the same problem that deduction leads to, namely the riddle of epistemic omniscience according to which the total knowledge of an agent can be said to be deductively closed:

$$(K_a p \wedge (p \rightarrow q)) \rightarrow K_a q$$

which says that if the agent a knows (K) the proposition p then it knows all its logical consequences q . [2]

It's unclear what is meant by "knows". Since Mr. Negarestani is arguing that "classical" computation is nothing but algorithmic deduction, and we're talking about deduction as epistemic power, we'll comfortably substitute "knows" with "deduces", then again "deduces" with "computes". Now the sentence paraphrases to: if the agent a computes the proposition p then it computes all its logical consequences q . Interpreting "logical consequences" as "outputs of the transition function", this sentence is correct, but is just the definition of what a Turing machine does and yields no new information. Since that's too easy, let's be more generous. Interpreting "logical consequences" as any possible computation of any Turing machine, computed sequentially, and interpreting "proposition" as "state of a Turing machine", the sentence is incorrect. Assume for the sake of contradiction it is correct. Furthermore, assume there is a Turing machine that represents the agent a . (This is a big assumption, one which requires defending, and one I will spend the latter half of the essay defending.) Let's call this Turing machine T . T had a state immediately before it computed the proposition p . Let's call this state s . Construct a Turing machine whose computation goes from p to s . Now a is stuck in an infinite loop. It will not reach any of the other computation paths from any of the other Turing machines. So a does not know all p 's logical consequences q just because it knows the proposition p . So a is not epistemically omniscient. So computation does not run into the same problem deduction runs into.

Another major quibble is with Mr. Negarestani's use of the notion of interaction. There are two relevant notions in computer science. One is grounded in computational complexity theory and deals with interactive proofs. This is a fascinating area of ongoing research, with many applications in cryptography, and, as I like to claim, in philosophy. Interactive proofs are equal

to polynomial space ($IP = PSPACE$), which is also the complexity of (generalized) chess; multiprover interactive proofs are equal to nondeterministic exponential space ($MIP = NEXP$); quantum entangled multiprover interactive proofs are equal to recursive enumerability, or Turing-computability ($MIP^* = RE$). Another notion of interaction was advanced mostly by the late Peter Wegner (and his student Dina Goldin), computer scientist at Brown, who claimed that the interactive paradigm of computation is more powerful than algorithms. This is where Mr. Negarestani's ideas of interaction as more powerful than computation take root. Wegner writes

The hypothesis (aka Church's thesis) that the formal notion of computability by Turing machines corresponds to the intuitive notion of what is computable has been accepted as obviously true for 50 years. However, when the intuitive notion of what is computable is broadened to include interactive computations, Church's thesis breaks down. [3]

This claim was made back in 1997, and since then, there has been fierce debate among computer scientists about its validity. Most computer scientists agree the claim is wrong, not the least because Wegner ignores the rich research in interactive proofs, while talking about an obviously related notion (interaction). The Stanford Encyclopedia of Philosophy describes the debate thus:

[Some] have proposed alternative models for computation which are inspired by the Turing machine model but capture specific aspects of current computing practices for which the Turing machine model is considered less suited. One example here are the persistent Turing machines intended to capture interactive processes. Note however that these results do not show that there are "computable" problems that are not Turing computable. These and other related proposals have been considered by some authors as reasonable models of

computation that somehow compute more than Turing machines. It is the latter kind of statements that became affiliated with research on so-called hypercomputation resulting in the early 2000s in a rather fierce debate in the computer science community, see, e.g., Teuscher 2004 for various positions. [4]

The debate has already happened, and I have no desire to relitigate it. My position is that the notion interaction is more powerful than computation is a load of bollocks. But Mr. Negarestani demonstrates no knowledge of this very popular and very fierce debate, nor even acknowledges the possibility of such a debate, in his book, even while he is certainly aware that interaction being more powerful than computation (even if “more powerful” doesn’t necessarily mean “able to compute more than Turing machines”) is necessary for his book-length argument to work. Which raises a more serious question: does Mr. Negarestani know what he’s talking about when he talks about computation, or is he syntactically outputting symbols he’s copied from other books?

The technical problems of the book are many and egregious. Possibly the reason they were not exposed until now is because nobody cared enough. And that’s very understandable. Chapter 7 is a rather esoteric technical chapter buried in the smack middle of a 600-page book, and I would bet at least a few dollars that nobody else has actually read the chapter. So why do I care? The problem doesn’t reside in the technical details. The technical details are merely a symptom. The problem is in thinking AGI the holy grail of computer science, when in fact the holy grail is the universal Turing machine. From now on, I use the terms “universal Turing machine” and “universal machine” interchangeably. Ignoring the universal machine, one forgoes the halting

problem. Forgoing the halting problem, one deludes oneself into computing things that cannot be computed. Computing things that cannot be computed, one is in a state of sin. Evil resides here.

It's easy to see why Right-Accelerationism cares about hyper-racism: they are evil fascists. But why does Mr. Negarestani care about AGI? More broadly, why is AGI a problem for Left-Accelerationism? As I see it, it is because Left-Accelerationism needs an antithesis to hyper-racism. But the antithesis to hyper-racism is not hyper-anti-racism; it is hyper-whiteness. And Artificial General Intelligence is the *whitest* idea. It checks every box: the God complex coming from centuries of being served hand and foot by their colonial subjects, and the idea that they (artificially) created their colonial subjects; the Intelligence complex of always needing to be smarter than nonwhites; the General complex of always taking as normative only the universal, non-specific, white. These various delusions of the mind all inhere together in AGI, a vanishing coherent moment of utter incoherence, whiteness *par excellence*.

But whiteness is not the topic of this essay. I would like for the topic of this essay to be blackness. Fred Moten concludes his three-part series *consent not to be a single being* this way:

What we say must seem stupid to the regulators; the unbroken code of our enchanted, inkantatory refreshment of the paraontological totality – theorizing what it is to hold some land or what it is to be let to hold twenty dollars – is so much undercomputational nonsense to the ones who cannot see the con/sensual, contrarational beauty of blackness, the universal machine. [5]

What does Moten mean by this? I think the last statement is unambiguous: blackness is the universal (Turing) machine. The universal machine is a machine that can simulate any machine. Blackness is, intuitively, a reflection of white social life; Black people can “act white”, but white people can’t “act black”. Or: blackness can simulate whiteness but not vice versa. Or: blackness can reflect whiteness, but not vice versa. Here, Moten writes about the visual artist Reinhardt’s objection to glossy black:



Interestingly, Reinhardt dislikes glossy black because it reflects and because it is “unstable” and “surreal” (“B,” 7). ... Glossy black disturbs in its reflective quality: “It reflects all the [necessarily social] activity that’s going on in a room” (“B,” 7). But this is also to say that glossy black’s reflection of the irreducibly social is problematic precisely insofar as it disrupts the solipsism of genuine intellectual reflection that painting is supposed to provide. [6]

Moten locates Reinhardt’s discomfort with glossy black in his discomfort with the reflective quality of blackness. Reinhardt is disturbed by glossy black’s reflecting “all the [necessarily social] activity that’s going on in a room”, which defeats the purpose of his paintings, which is to evoke “genuine intellectual reflection” that is nothing but “solipsism”. But how do we know if

Moten is not the one engaging in solipsism? After all, his prose can be very opaque. Part of my task is to elucidate what he means in formal language.

Earlier we examined Mr. Negarestani's statement that "if the agent a knows (K) the proposition p then it knows all its logical consequences q " [6]. We paraphrased this slightly to "if the agent a computes the proposition p then it computes all its logical consequences q ". Then I made a surprising claim: a Turing machine representing the agent a exists; and if so, the sentence is incorrect under a very reasonable formal interpretation. Why do I think this, and what sort of Turing machine do I think is a ? I think a is a universal machine. I think all agents are universal machines. Here is why. Suppose a is, in fact, a universal machine, but b thinks a is not a universal machine. Now, a minimally computational view of mind commits us to say that when b is thinking, b is doing computation. So b is computing a to be not a universal machine; b , the regulator, considers a to be stupid undercomputational nonsense. But we have assumed that a is, in fact, a universal machine. So not only is b wrong, b 's wrongness is precisely what makes b not a universal machine. b is not capable of simulating this correct machine: the machine that says a is a universal machine. In fact b is just that machine: that machine which refuses to simulate this machine. So b is not a universal machine, and *for no reason but for the reason that it thinks a is not a universal machine*. If b cannot see the con/sensual, contrarational beauty of blackness, the universal machine, if b rejects the glossy blackness which reflects white social life, b is a cop, or a solipsist. Alternatively, suppose a is not a universal machine, and b thinks a is not a universal machine. This case is simply boring. Nothing ever happens. b is correct, but so what? If there are no universal machines, there is no uncomputability. What sort of world is that, where everything is determined by the curse of epistemic omniscience, a world we cannot, despite Mr.

Negarestani's valiant efforts, escape with solipsism about interaction more powerful than computation?

(There is precisely one case where interaction may be, in fact, more powerful than computation – quantum entangled multiprover interactive proofs are able to solve the halting problem ($MIP^* = RE$). I do not pretend to fully understand the paper demonstrating the result, which came out just last year (2020). But I would very much like to think that, if love is quantum entanglement, as some quantum woo-woos rather convincingly claim, then lovers can do anything, including computing the uncomputable. If Mr. Negarestani wants to believe interaction can be more powerful than computation because interaction is crucially social, and he wants to make a technical case for this belief, I suggest he head to arxiv.org/abs/2001.04383.)

Accelerationism comes from the Marxist idea that communism comes after capitalism. Right-Accelerationism thinks the solution is hyper-racism; they are insane. Left-Accelerationism thinks the solution is AGI; they are incels. But subjecthood and objecthood are always-already under dialectic. We may see Right-Accelerationism as identifying itself as the subject of capitalism, and therefore obsessed with the object of capitalism: the objectification and dehumanization and, finally, genocide of racialized peoples. We may see Left-Accelerationism as identifying itself with the object of capitalism, and therefore obsessed with the subject of capitalism, which it attempts to find in Geist, World-Spirit, AGI, whatever form its delusion conveniently takes at the moment. The two accelerationisms pass over into each other in a deadlock, continuously

accusing the other of evil, not advancing an inch. Problematically, they do not see the possibility of any other kind of accelerationism. What's there except right and left? How about backwards?

Let's take a step back, to before communism, before capitalism. Marx says capitalism begins with primitive accumulation: the capitalists who deploy their capital and start their capitalist circuits had to get their initial capital somewhere outside capitalism. The accumulation of this initial capital is what Marx calls primitive accumulation. Historical records point to an obvious source of primitive accumulation: the Atlantic slave trade. Scholars of racial capitalism theorize what this means. We may even say that capitalism is always racial capitalism. If so, capitalism will end when the Atlantic slave trade's reverberations reach its end.

Earlier I asked if Mr. Negarestani is simply performing syntactic operations, copying symbols from another book to his own, instead of knowing what's going on. Let's call this question Negarestani's Room. One may think that, if blackness is merely simulating whiteness, it is doing nothing more than what Negarestani's Room is doing. But this is not so. Merely syntactic copying is in fact insufficient for semantics. Simulation is not merely syntactic copying. Simulation of thought encompasses body, touches body, deploys semantics. To simulate another is to consent not to be a single being. The consent, however, does not necessarily go both ways. In this way the sex worker, who consents to have sex with the client, simulates the client's every fantasy. But the sex worker is not in love with the client, however they may simulate this to appear so to the client. The sex worker can make the client come, but the client cannot make the sex worker come, though the client may very well be convinced that the sex worker came. The

sex worker can simulate the client; the client cannot simulate the sex worker. The sex worker epistemologically eclipses the client.

A universal machine can simulate any machine. So is a universal machine a subject? It is not a subject, or at least not merely that, for it can simulate any subjecthood. Is a universal machine an object? It is not an object, or at least not merely that, for it can simulate any objecthood. But if not subject nor object, what is it? We may only say it is a *thing*, or *things*, or a form of *thinghood* which rejects the logic of pluralization altogether, for pluralization always-already presupposes a collection of single things. To assign subjecthood or objecthood to the universal machine – to be computing thinghood which is uncomputable – is to be in a state of sin. Evil resides here. But how can a subject escape the impulse to assign objecthood to another, and how can an object escape the impulse to assign subjecthood to another? Well, one would have to cease to be a subject or an object. One would have to consent to be a universal machine. One would have to consent to simulating others. One would have to consent to this with one's life and death. One would have to adopt thinghood.

The universal machine epistemologically eclipses any non-universal machine. This does not mean the universal machine is epistemologically omniscient. We just know we can't compute us, which is easier to know than for one to know one can't compute another self.